

Assignment Project Exam Help

COMP90015 Distributed Systems
Socket Model and Threading Paradigms

<https://eduassistpro.github.io>

School of Computing and Informati

© The University of Melb

Add WeChat edu_assist_pr
2022 Semester II

Assignment Project Exam Help

1 Socket Model


2 Queue

<https://eduassistpro.github.io>

3 Threading Paradigms

Add WeChat edu_assist_pro

Essential aspects I



The essential elements of modeling a socket include:

- *incoming socket connections* – communication requests, e.g., using the TCP protocol, that are requesting that a socket connection be established.
- The client initiating the request will time out if the request is lost, e.g. by network errors or if the OS of the server drops the request due to lack of resources, and possibly retry several times. This can be a significant source of delay when using connection oriented protocols like TCP over an unreliable network.

Essential aspects II

- Loss of the request in the network is more of a problem than the OS dropping the request since the OS can at least signal the sender that the request was dropped, whereas in the former case the sender will wait for the entire timeout period before it deems the request to be lost.
- *OS queue* – If the socket connection is valid, i.e. to a port that a process is bound to, then the socket connection is put into an OS queue.
 - The OS maintains a queue of OS threads.
 - If the queue is full, the connection is dropped (lost). The OS may signal the receiver that the connection is refused or dropped in this case.
 - Some OSes have a process like the UNIX `xinet` that listens on many of the well-known port numbers and starts the relevant server processes listening on that port. This is preferable to having many different processes started in advance, in the case when connections are infrequent, since they represent OS overheads.
 - Some OSes let multiple processes accept connections on the same port. All such processes simply make use of the same OS queue.
- *IO thread* – If a process binds to a port for socket based communication then it needs to have at least one thread that accepts socket connections and processes them, e.g. accessing local data, reading/writing to the socket.

Essential aspects III

Assignment Project Exam Help

- At any one time the IO thread is either processing a socket connection, or is blocked (idle) while it is waiting for a new connection to arrive.
- The IO thread will not return even if a socket was not waiting on the queue.
- If the thread does not process incoming socket requests fast enough, requests will start dropping at the OS i.e. they will be lost, and the client will have to try the request later.

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro

Queueing theory

The Socket Model can be described using queueing theory:

- The socket connection request arrives at a mean rate of λ requests per second, and we can use an Exponential distribution with parameter λ , similar to failure rate models, to describe the probability of a request arriving within the next t seconds, which we recall is the cumulative distribution function of the Exponential distribution.
- The time a request spends in the queue (the time t from the time the request arrives at the queue to the time it is processed), can be modeled as well using an Exponential distribution with parameter μ , called the service rate, i.e. the mean number of requests that the IO thread can process per second.
- If the queue has a finite capacity of $k - 1$ so that at most k requests can be in the system (further 1 socket connection possibly being currently processed by the IO thread), making a maximum of k socket connections in the system at any one time, then the essential aspects of the Socket Model are described using a $M/M/1/k$ queue.

Queueing theory results I

https://sites.pitt.edu/~dttipper/2130/2130_Slides4.pdf

Without looking at deriving queueing theory results, some of the most relevant results are:

- Since the queue has finite capacity $k - 1$, it can sometimes become full and new socket requests received by the OS will have to be dropped. The drop rate is equivalently the blocking probability of the queue:

<https://eduassistpro.github.io>

- For constant $\rho < 1$, if $k \rightarrow \infty$ (unbounded queue) t
- The portion of socket requests dropped is $\lambda - \lambda_{\text{eff}}$. The effective socket request rate that the IO thread sees as λ_{eff} .
- The effective thread utilization is then $\frac{\lambda_{\text{eff}}}{\mu}$, which is the fraction of time that the thread will be busy, rather than idle (waiting for a request to arrive), and the system is said to be *stable*.
 - When $\lambda_{\text{eff}} \geq \mu$ then the thread will eventually be busy 100% of the time, meaning it will eventually fall behind, and the system is said to be *unstable*.

Queueing theory results II

https://sites.pitt.edu/~dttipper/2130/2130_Slides4.pdf

- The average number of socket requests in the system becomes:

$$L = \frac{\rho}{1-\rho} - \frac{(k+1)\rho^{k+1}}{1-\rho^{k+1}} \quad \rho \neq 1$$

- The L is the sy and so the queue length becomes $L_q = L - \rho_{eff}$.


- For constant $\rho < 1$, if $k \rightarrow \infty$ then this becomes —

- Notice that in either case the mean queue length sharply increases a vertical asymptote. This is why we prefer using a bounded queue consumption can sharply increase under high load *capacity of the server*.

- The average time that a socket request spends in the system, including the time taken by the IO thread to process the socket request (i.e. to close the socket), is $W = \frac{L}{\lambda_{eff}}$.

- The average time that a socket request spends waiting in the queue is $W - \frac{1}{\mu}$.

Thread-per-connection I



The thread-per-connection paradigm creates a new thread for every socket connection:

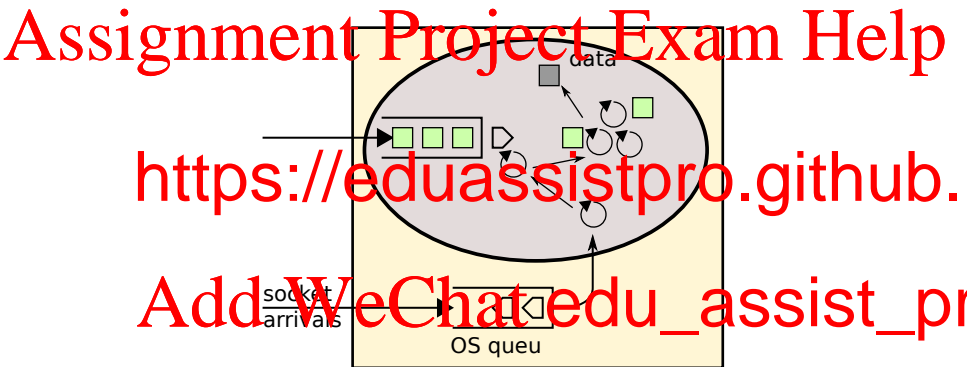
- Creating more threads is one way to utilize more cores of a multi-core server.

Thread-per-connection II

- If the maximum number of threads allowed to be created by the IO thread is c then the system is described by an $M/M/c/k$ queue, which will provide different results to what we have seen so far. In this case the IO thread is not considered to be doing any useful work, i.e. it never actually processes a socket connection itself.
- As the total number of threads increases beyond the number of processors (i.e. c), the system becomes more congested (i.e. the queue length increases as well). Context switching overheads also increase as the number of threads increases. Context state needs to be stored in cache and pushes useful data out of the cache. Context switching is also more expensive as the number of threads increases.
- Multiple threads leads to concurrency control requirements. Concurrency control leads to further context switching and synchronization overheads between threads.
- Lock-free designs can greatly reduce the amount of context switching required, especially by using special machine instructions, however lock-free designs cannot block and therefore data in a distributed system must be dropped when the distributed systems' resources are exceeded, which is not necessarily desired.

Thread-per-request I

Depicted combined with thread-per-connection



The thread-per-request paradigm goes further and allows the thread that is processing the socket to create threads for each request received on the socket:

Thread-per-request II

Depicted combined with thread-per-connection

Assignment Project Exam Help

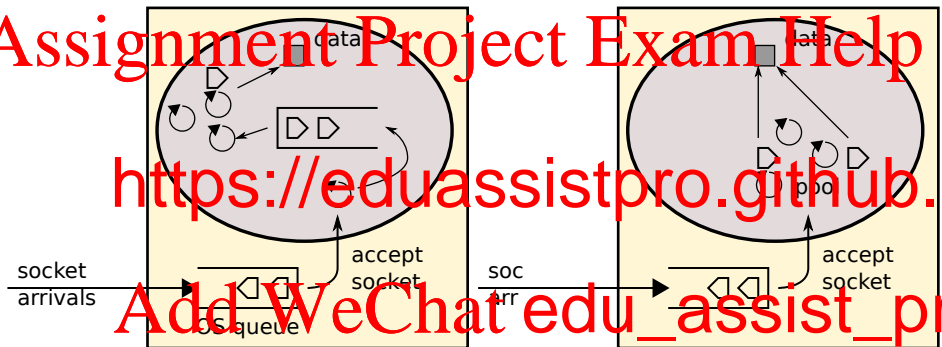
- Even more threads are potentially created than the thread-per-connection paradigm
- Some per-request threads are also created, so thread-per-request is not strictly thread-per-connection.
- Requests may have dependencies between them, in that the order that the requests are processed may be important.

- If many requests on the socket can be bundled into one large item on the queue, then this can reduce context switching at a rate proportional to the size of the bundles.

<https://eduassistpro.github.io>


Add WeChat [edu_assist_pro](https://eduassistpro.github.io)

Thread pool



The thread-pool paradigm creates a fixed or dynamically resizable set of threads that either take incoming socket connection requests from a process maintained queue, or, if the accept socket API is thread safe then the pool of threads can take directly from the OS queue.

Thread-per-object paradigm



Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pr

The thread-per-object paradigm creates a thread for each data object. This can potentially reduce cache miss rates in the machine since for a given data object, only 1 thread ever accesses it and it will reside only in the cache for that thread. This can greatly increase cache efficiency which can significantly improve overall performance of the machine.

Discussion Question

Assignment Project Exam Help

Question

better than
processes

`-process` is

ed
threads.

Critically compare the two approaches and discuss what are the main aspects of your comparison.

Add WeChat edu_assist_pr